

Exposing pods to the cluster

Create an nginx Pod, and note that it has a container port specification:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-nginx
spec:
  selector:
    matchLabels:
      run: my-nginx
  replicas: 2
  template:
    metadata:
      labels:
        run: my-nginx
    spec:
      containers:
        - name: my-nginx
          image: nginx
          ports:
            - containerPort: 80
```

This makes it accessible from any node in your cluster. Check the nodes the Pod is running on:

```
$ kubectl create -f ./run-my-nginx.yaml
```

```
$ kubectl get pods -l run=my-nginx -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
my-nginx-3800858182-jr4a2	1/1	Running	0	13s	10.244.3.4	kubernetes-minion-905m
my-nginx-3800858182-kna2y	1/1	Running	0	13s	10.244.2.5	kubernetes-minion-ljyd

Check your pods' IPs:

```
$ kubectl get pods -l run=my-nginx -o yaml | grep podIP
  podIP: 10.244.3.4
  podIP: 10.244.2.5
```

Creating a Service

So we have pods running nginx in a flat, cluster wide, address space. In theory, you could talk to these pods directly, but what happens when a node dies? The pods die with it, and the Deployment will create new ones, with different IPs. This is the problem a Service solves.

You can create a Service for your 2 nginx replicas with **kubectl expose**:

```
$ kubectl expose deployment/my-nginx
service/my-nginx exposed
```

This is equivalent to `kubectl create -f` the following yaml:

```
apiVersion: v1
kind: Service
metadata:
  name: my-nginx
  labels:
    run: my-nginx
spec:
  ports:
    - port: 80
      protocol: TCP
  selector:
    run: my-nginx
```

[Concepts](#)

Overview

[What is Kubernetes?](#)

[Kubernetes Components](#)

[The Kubernetes API](#)

Working with Kubernetes Objects

[Understanding Kubernetes Objects](#)

[Names](#)

[Namespaces](#)

[Labels and Selectors](#)

[Annotations](#)

[Field Selectors](#)

[Recommended Labels](#)

Object Management Using kubectl

[Kubernetes Object Management](#)

[Managing Kubernetes Objects Using Imperative Commands](#)

[Imperative Management of Kubernetes Objects Using Configuration Files](#)

[Declarative Management of Kubernetes Objects Using Configuration Files](#)

Kubernetes Architecture

[Nodes](#)

[Master-Node communication](#)

[Concepts Underlying the Cloud Controller Manager](#)

Containers

[Images](#)

[Container Environment Variables](#)

[Runtime Class](#)

[Container Lifecycle Hooks](#)

Workloads

Pods

[Pod Overview](#)

[Pods](#)

[Pod Lifecycle](#)

[Init Containers](#)

[Pod Preset](#)

[Disruptions](#)

Controllers

[ReplicaSet](#)

[ReplicationController](#)

[Deployments](#)

[StatefulSets](#)

[DaemonSet](#)

[Garbage Collection](#)

[TTL Controller for Finished Resources](#)

[Jobs - Run to Completion](#)

[CronJob](#)

Services, Load Balancing, and Networking

[Services](#)

[DNS for Services and Pods](#)

[Connecting Applications with Services](#)

[Ingress](#)

[Network Policies](#)

[Adding entries to Pod /etc/hosts with HostAliases](#)

Storage

[Volumes](#)

[Persistent Volumes](#)

[Volume Snapshots](#)

[Storage Classes](#)

[Volume Snapshot Classes](#)

[Dynamic Volume Provisioning](#)

[Node-specific Volume Limits](#)

Configuration

[Configuration Best Practices](#)

[Managing Compute Resources for Containers](#)

[Assigning Pods to Nodes](#)

[Taints and Tolerations](#)

[Secrets](#)

[Organizing Cluster Access Using kubeconfig Files](#)

[Pod Priority and Preemption](#)

[Scheduler Performance Tuning](#)

Policies

[Resource Quotas](#)

[Pod Security Policies](#)

Cluster Administration

[Cluster Administration Overview](#)

[Certificates](#)

[Cloud Providers](#)

[Managing Resources](#)

[Cluster Networking](#)

[Logging Architecture](#)

[Configuring kubelet Garbage Collection](#)

[Federation](#)

[Proxies in Kubernetes](#)

[Controller manager metrics](#)

[Installing Addons](#)

Extending Kubernetes

[Extending your Kubernetes Cluster](#)

Extending the Kubernetes API

[Extending the Kubernetes API with the aggregation layer](#)

[Custom Resources](#)

Compute, Storage, and Networking Extensions

[Network Plugins](#)

[Device Plugins](#)

[Service Catalog](#)

[Edit This Page](#)

Connecting Applications with Services

The Kubernetes model for connecting containers

Now that you have a continuously running, replicated application you can expose it on a network. Before discussing the Kubernetes approach to networking, it is worthwhile to contrast it with the “normal” way networking works with Docker.

By default, Docker uses host-private networking, so containers can talk to other containers only if they are on the same machine. In order for Docker containers to communicate across nodes, there must be allocated ports on the machine’s own IP address, which are then forwarded or proxied to the containers. This obviously means that containers must either coordinate which ports they use very carefully or ports must be allocated dynamically.

Coordinating ports across multiple developers is very difficult to do at scale and exposes users to cluster-level issues outside of their control. Kubernetes assumes that pods can communicate with other pods, regardless of which host they land on. We give every pod its own cluster-private-IP address so you do not need to explicitly create links between pods or mapping container ports to host ports. This means that containers within a Pod can all reach each other’s ports on localhost, and all pods in a cluster can see each other without NAT. The rest of this document will elaborate on how you can run reliable services on such a networking model.

This guide uses a simple nginx server to demonstrate proof of concept. The same principles are embodied in a more complete [Jenkins CI application](#).

- [Exposing pods to the cluster](#)
- [Creating a Service](#)
- [Accessing the Service](#)
- [Securing the Service](#)
- [Exposing the Service](#)
- [What's next](#)

Exposing pods to the cluster

We did this in a previous example, but let’s do it once again and focus on the networking perspective. Create an nginx Pod, and note that it has a container port specification:

[service/networking/run-my-nginx.yaml](#)




```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-nginx
spec:
  selector:
    matchLabels:
      run: my-nginx
  replicas: 2
  template:
    metadata:
      labels:
        run: my-nginx
    spec:
      containers:
        - name: my-nginx
          image: nginx
          ports:
            - containerPort: 80

```

This makes it accessible from any node in your cluster. Check the nodes the Pod is running on:

```
$ kubectl create -f ./run-my-nginx.yaml
```

```
$ kubectl get pods -l run=my-nginx -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
my-nginx-3800858182-jr4a2	1/1	Running	0	13s	10.244.3.4	kubernetes-minion-905m
my-nginx-3800858182-kna2y	1/1	Running	0	13s	10.244.2.5	kubernetes-minion-ljyd

Check your pods' IPs:

```
$ kubectl get pods -l run=my-nginx -o yaml | grep podIP
```

```
  podIP: 10.244.3.4
```

```
  podIP: 10.244.2.5
```

You should be able to ssh into any node in your cluster and curl both IPs. Note that the containers are *not* using port 80 on the node, nor are there any special NAT rules to route traffic to the pod. This means you can run multiple nginx pods on the same node all using the same containerPort and access them from any other pod or node in your cluster using IP. Like Docker, ports can still be published to the host node's interfaces, but the need for this is radically diminished because of the networking model.

You can read more about [how we achieve this](#) if you're curious.

Creating a Service

So we have pods running nginx in a flat, cluster wide, address space. In theory, you could talk to these pods directly, but what happens when a node dies? The pods die with it, and the Deployment will create new ones, with different IPs. This is the problem a Service solves.

A Kubernetes Service is an abstraction which defines a logical set of Pods running somewhere in your cluster, that all provide the same functionality. When created, each Service is assigned a unique IP address (also called clusterIP). This address is tied to the lifespan of the Service, and will not change while the Service is alive. Pods can be configured to talk to the Service, and know that communication to the Service will be automatically load-balanced out to some pod that is a member of the Service.

You can create a Service for your 2 nginx replicas with kubectl expose:

```
$ kubectl expose deployment/my-nginx
service/my-nginx exposed
```

This is equivalent to kubectl create -f the following yaml:

[service/networking/nginx-svc.yaml](#)



```
apiVersion: v1
kind: Service
metadata:
  name: my-nginx
  labels:
    run: my-nginx
spec:
  ports:
    - port: 80
      protocol: TCP
  selector:
    run: my-nginx
```

This specification will create a Service which targets TCP port 80 on any Pod with the run: my-nginx label, and expose it on an abstracted Service port (targetPort: is the port the container accepts traffic on, port: is the abstracted Service port, which can be any port other pods use to access the Service).

```
$ kubectl get svc my-nginx
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
my-nginx	ClusterIP	10.0.162.149	<none>	80/TCP	21s

As mentioned previously, a Service is backed by a group of Pods. These Pods are exposed through endpoints. The Service's selector will be evaluated continuously and the results will be POSTed to an Endpoints object also named my-nginx. When a Pod dies, it is automatically removed from the endpoints, and new Pods matching the Service's selector will automatically get added to the endpoints. Check the endpoints, and note that the IPs are the same as the Pods created in the first step:

```
$ kubectl describe svc my-nginx
```

```
Name:          my-nginx
Namespace:     default
Labels:        run=my-nginx
Annotations:    <none>
Selector:      run=my-nginx
Type:          ClusterIP
IP:            10.0.162.149
Port:          <unset> 80/TCP
Endpoints:     10.244.2.5:80,10.244.3.4:80
Session Affinity: None
Events:        <none>
```

```
$ kubectl get ep my-nginx
```

NAME	ENDPOINTS	AGE
my-nginx	10.244.2.5:80,10.244.3.4:80	1m

Accessing the Service

Kubernetes supports 2 primary modes of finding a Service - environment variables and DNS. The former works out of the box while the latter requires the [CoreDNS cluster addon](#).

Environment Variables

When a Pod runs on a Node, the kubelet adds a set of environment variables for each active Service. This introduces an ordering problem. To see why, inspect the environment of your running nginx Pods (your Pod name will be different):

```
$ kubectl exec my-nginx-3800858182-jr4a2 -- printenv | grep SERVICE
KUBERNETES_SERVICE_HOST=10.0.0.1
KUBERNETES_SERVICE_PORT=443
KUBERNETES_SERVICE_PORT_HTTPS=443
```

Note there's no mention of your Service. This is because you created the replicas before the Service. Another disadvantage of doing this is that the scheduler might put both Pods on the same machine, which will take your entire Service down if it dies. We can do this the right way by killing the 2 Pods and waiting for the Deployment to recreate them. This time around the Service exists *before* the replicas. This will give you scheduler-level Service spreading of your Pods (provided all your nodes have equal capacity), as well as the right environment variables:

```
$ kubectl scale deployment my-nginx --replicas=0; kubectl scale deployment my-nginx --replicas=2;
```

```
$ kubectl get pods -l run=my-nginx -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
my-nginx-3800858182-e9ihh	1/1	Running	0	5s	10.244.2.7	kubernetes-minion-ljyd
my-nginx-3800858182-j4rm4	1/1	Running	0	5s	10.244.3.8	kubernetes-minion-905m

You may notice that the pods have different names, since they are killed and recreated.

```
$ kubectl exec my-nginx-3800858182-e9ihh -- printenv | grep SERVICE
KUBERNETES_SERVICE_PORT=443
MY_NGINX_SERVICE_HOST=10.0.162.149
KUBERNETES_SERVICE_HOST=10.0.0.1
MY_NGINX_SERVICE_PORT=80
KUBERNETES_SERVICE_PORT_HTTPS=443
```

DNS

Kubernetes offers a DNS cluster addon Service that automatically assigns dns names to other Services. You can check if it's running on your cluster:

```
$ kubectl get services kube-dns --namespace=kube-system
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kube-dns	ClusterIP	10.0.0.10	<none>	53/UDP,53/TCP	8m

If it isn't running, you can [enable it](#). The rest of this section will assume you have a Service with a long lived IP (my-nginx), and a DNS server that has assigned a name to that IP (the CoreDNS cluster addon), so you can talk to the Service from any pod in your cluster using standard methods (e.g. gethostbyname). Let's run another curl application to test this:

```
$ kubectl run curl --image=radial/busyboxplus:curl -i --tty
```

Waiting for pod default/curl-131556218-9fnch to be running, status is Pending, pod ready: false

Hit enter for command prompt

Then, hit enter and run nslookup my-nginx:

```
[ root@curl-131556218-9fnch:/ ]$ nslookup my-nginx
```

```
Server: 10.0.0.10
Address 1: 10.0.0.10
```

```
Name: my-nginx
Address 1: 10.0.162.149
```

Securing the Service

Till now we have only accessed the nginx server from within the cluster. Before exposing the Service to the internet, you want to make sure the communication channel is secure. For this, you will need:

- Self signed certificates for https (unless you already have an identity certificate)
- An nginx server configured to use the certificates
- A [secret](#) that makes the certificates accessible to pods

You can acquire all these from the [nginx https example](#). This requires having go and make tools installed. If you don't want to install those, then follow the manual steps later. In short:

```
$ make keys secret KEY=/tmp/nginx.key CERT=/tmp/nginx.crt SECRET=/tmp/secret.json
```

```
$ kubectl create -f /tmp/secret.json
```

```
secret/nginxsecret created
```

```
$ kubectl get secrets
```

NAME	TYPE	DATA	AGE
default-token-il9rc	kubernetes.io/service-account-token	1	1d
nginxsecret	Opaque	2	1m

Following are the manual steps to follow in case you run into problems running make (on windows for example):

#create a public private key pair

openssl req -x509 -nodes -days 365 -newkey rsa:2048 -keyout /d/tmp/nginx.key -out /d/tmp/nginx.crt -subj "/CN=my-nginx/O=my-nginx"

#convert the keys to base64 encoding

cat /d/tmp/nginx.crt | base64

cat /d/tmp/nginx.key | base64

Use the output from the previous commands to create a yaml file as follows. The base64 encoded value should all be on a single line.

apiVersion: "v1"

kind: "Secret"

metadata:

name: "nginxsecret"

namespace: "default"

data:

nginx.crt:

"LS0tLS1CRUdJTiBDRVJUSUZJQ0FURSB0tLS0tCk1JSURlekNDQWdlZ0F3SUJBZ0lKQUp5M3lQK0pzMlplJTUEwR0NTcUdTSWlzMlRFRFFkRVRVFNQ1I4RVRBUEJnTIYKQkFNVENHNW5hVzU0YzNaak1SRXdEd1IEVIFRS0V3aHVaMmx1ZUhOMlI6QWVGdzB4TnpFd01qWXdOekEzTVRKYQpGdzB4T0RFd01qWXdOekEzTVRKYU1DWXhFVEFQQmdOVkJBTVRDRzVuYVc1NGMzMmpNUkV3RHdZRFZRUUtdF2h1CloybHVISE4yWXPdQ0FTSXdEUVIKS29aSWWh2Y05BUUVVCQIFBRGdnRVBBRENDQVFvQ2dnRUJBSjFjSU1SOVdWM0IKMIZIQIRMRmtobDRONXlJMEJxYUhlQktMSnJMcy8vdzZhU3hRS29GbHJSU94NGUrMIN5ajBFcndCLzIYTnBwbQppeW1CL3JkRldkOXg5UWWhBQUxkZkVhTmNiV3NsTVFVcnhBZW50VWt1dk1vLzgvMHRpbGhjc3paenJEYVJ4NEo5Ci82UVRtVVI3a0ZTWUpOWTVQZkR3cGc3dlVvaDZmZ1Voa m92VG42eHNVR0M2QURVODBPnXFIZWhNeVI1N2lmU2YKNHZpaXdIY3hnL3lZR1JBRS9mRTRqakxCdmdONjc2SU90S01rZXV3R0ljNDFhd05tNnNTSszRqYUNGEGpYSnZaZQp2by9kTIEybHhHWCtKT2l3SEhXbXNhGp4WTRaNVk3R1ZoK0QrWnYvcW1mMFgvcVY0Rmo1NzV3ajFMWVBocWtsCmdhSXZYRyt4U1FVQ0F3RUFBU5RTUU0d0hRWURWUjBPQkJZRUZPNG9OWkl3YXc1OUlsYkROMzhIYkduYnhFVjcKTUI4R0ExVWRJd1FZTUJhQUZPNG9OWkl3YXc1OUlsYkROMzhIYkduYnhFVjdNQXdhQTFVZEV3UUZnQU1CQWY4dWpEUVIKS29aSWWh2Y05BUUVGQIFBRGdnRUJBRVhTMW9FU0lFaXdyMDhWcVA0K2NwTHI3TW5FMTducDBvMm14aFVjRGb0RvRjdRZnZqeE04Tzd2TjB0clcx2pGSW0vWDE4ZnZaL3k4ZzVaWG40Vm8zc3hKVmRBcStNZC9jTStzUGEKNmJjTkNUekZqeFpUV0UrKzE5NS9zb2dmOUZ3VDVDK3U2Q3B5N0M3MTZvUXRUakViV05VdEt4cXI0Nk1OZWNCMApwRFhWZmdWQTRadkR4NFo3S2RiZDY5eXM3OVFHYmg5ZW1PZ05NZFIsSUswSGt0ejF5WU4vbVpmK3FqTkJqbWZjCkNnMnlwbGQ0Wi8rUUNQZjl3SkoybFlrY2FnT0R4eIBWcGxNSEcybzgvTHFDdnh6elZPUDUxeXdLZEtaUMwSVEKQ0l5T2wwWW5scE9UNEh1b2hSUzBPOSTlMm9KdFZsNUlyczRpbDIhZ3RTVXFxUIU9Ci0tLS0tRU5EIENFUlRJRkdQVRFLS0tLS0K"

nginx.key:

"LS0tLS1CRUdJTiBQUkIWQVRFIETFWs0tLS0tCk1JSUV2UUICQURBTkNa3Foa2IHOXcwQkFRRUZBQVNDQktjd2dnU2pBZ0VBQW9JQkFRQ2RhaURFZlZsZHdkbFIkd1V5eFpJWmVEZWNUtkFhbWh4d1NpeWF5N1AvOE9ta3NVQ3FCWmNpQ0RzZUh2dGtzbzICSzhBZi9WemFhWm9zcApnZjYzUIZuZmNmVUIRQUUN3WHhHVfHhMXJKVEVGSzhRSHA3VkpMcnpLUC9QOUxZcFIYTE0yYzZ3MmtjZUNmZitrCkU1bEVINUJVbUNUV09UM3c4S1IPNzFLSWVuNEZJWTZMMDUrc2JGQmd1Z0ExUE5JdWFubm9UTWtIZTRuMG4rTDQKb3NCM01ZUDhtQmtRQIAzeE9JNHI3YjREZXUraURyU2pKSHJzQmIIT05Xc0RadXJFaXVJMmdoY1kxeWlyWHI2UAozVFVOCnSbC9pVG9zQngxcHJHclK4V09HZVdPeGxZZmcbWlVnNBUOUYvNWxlQlkrZStjSTITMkQ0YXBKWUdpCkwxeHZzVWtGQWdNQkFBRUNnZ0VBZFhCK0xkbk8ySEIOTGo5bWRsb25IUGIHWWVzZ294RGQwci9hQ1Zkank4dIEKTjlwL3FQWkUxek1yall6Ry9kVGhTMMmWc0QxaTBXSjdwr1IGb0xtdXIWTjltY0FXUTM5SjM0VHZaU2FFSWZWNGo5TE1jUHHNTmFsNjRLMFRVbUFQZytGam9QSFIhUuXLOERLOUtnNXNrSE5pOWNzMIY5ckd6VWIVZWtBL0RBUIBTCII3L2ZjUFBacDRuRWVBZml3WTK1R1Ilb1p5V21SU3VKdINyblBESGtUdW1vVIVWdkxMRHRzaG9reUxiTWVtN3oKMmJzVmpwSW1GTHJqbGtmQXlpNHg0WjJrV3YyMFRrdWtsZU1jaVIMbjk4QWxiRi9DSmRLM3QraTRoMTVIR2ZQegpoTnh3bk9QdIVTaDR2Q0o3c2Q5TmtEUGJvS2JneVVHOXBYamZhRGR2UVFLQmdRRFFLM01nUkhkQ1pKNVFqZWFKCIFGdXF4cHdnNzhZTjQyL1NwenIUYmtGcVFoQWtyczJxWGx1MDZBRzhrZzlzQkswaHkzaE9zSGgxcXRVK3NHZVAKOWRERHBsUWV0ODZsY2FIR3hoc0V0L1R6cEdtNGFKSm5oNzVVaTVGZk9QTDhPTm1FZ3MxMVRhUldhNzZxeIRyMgphRlpjQ2pWV1g0YnRSTHVwSkgrMjZnY0FhUUtCZ1FEQmxVSUuzTnNVOFBZYEVL25sQVB5VWs1T3IDdWc3dmVyCIUycXlrdXFzYnBkSi9hODViT1JhM05IVmpVM25uRGpHVHBWaeE9JeXg5TEFrc2RwZEFjVmxvcG9HODhXYk9IMTAKMUdqbnkySmdDK3JVWUZiRGtpUGx1K09IYnRnOXFYcGJMSHBzUVpsMGhucDBYSFNyVm9CMUliQndnMGEyOFVadApCbFBtWmc2d1BRS0JnRHVIUVV2SDZHYTNDVUsxNFdmOFhlcFFnMU16M2VvWTBPQm5iSDRvZUZKZmcraEppSXlnCm9RN3hqWldVR3Blc3AyblRtcHERQWISNzdyRVhsdlhtOEIVU2FsbkNiRGIKY01Pc29RdFBZNS9NczJMRm5LQTQKaENmL0pWb2FtZm1nZEN0ZGtFMXNINE9MR2IJVHdEbTRpb0dWZGlwMlInbzFyb2htNUpLMUI3MkpBb0dBuW01UQpHNDhXOTVhL0w1eSt5dCsyZ3YvUHM2VnBvMjZITzRNQ3IJazJVem9ZWE9IYnNkODJkaC8xT2sybGdHZlI2K3VuCnc1YytZUXRSTHhQmd3MUtpbGhFZDBKTWU3cGpUSVpnQWJ0LzVPbnIDak9OVXN2aDJJS2lrQ1Z2dTZsZlBjNkQKckliT2ZlaHhxV0RZK2Q1TGN1YSst2NzJ0RkxhenJsSIBsRzIOZHhrQ2dZRUF5ellzT3UyMDNRVvV6bUICRkwzZAp4Wm5XZ0JLSEo3TnNxcGFwb2RjL0d5aGVycjFDZzE2MmJaSjJDV2RsZkl0VEdtUjZZdmxTZEFOOFrWUWhFbUtKCnFBLzVzdHdxNWd0WGVLOVJmMWxXK29xNThRNTBxBmk1NVdUTThoSDZhTjlaMTItZ0FGdE5VdGNqQUx2dFYxdEYKWSS4WFJkSHJaRnBIWII2NWkwVW1VbGc9Ci0tLS0tRU5EIFBSSVZBVEUgS0VZLS0tLS0K"

Now create the secrets using the file:

```
$ kubectl create -f nginxsecrets.yaml
```

```
$ kubectl get secrets
```

NAME	TYPE	DATA	AGE
default-token-il9rc	kubernetes.io/service-account-token	1	1d
nginxsecret	Opaque	2	1m

Now modify your nginx replicas to start an https server using the certificate in the secret, and the Service, to expose both ports (80 and 443):

```
apiVersion: v1
kind: Service
metadata:
  name: my-nginx
  labels:
    run: my-nginx
spec:
  type: NodePort
  ports:
    - port: 8080
      targetPort: 80
      protocol: TCP
      name: http
    - port: 443
      protocol: TCP
      name: https
  selector:
    run: my-nginx
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-nginx
spec:
  selector:
    matchLabels:
      run: my-nginx
  replicas: 1
  template:
    metadata:
      labels:
        run: my-nginx
    spec:
      volumes:
        - name: secret-volume
          secret:
            secretName: nginxsecret
      containers:
```



```
- name: nginxhttps
  image: bprashanth/nginxhttps:1.0
  ports:
    - containerPort: 443
    - containerPort: 80
  volumeMounts:
    - mountPath: /etc/nginx/ssl
      name: secret-volume
```

Noteworthy points about the nginx-secure-app manifest:

- It contains both Deployment and Service specification in the same file.
- The [nginx server](#) serves HTTP traffic on port 80 and HTTPS traffic on 443, and nginx Service exposes both ports.
- Each container has access to the keys through a volume mounted at /etc/nginx/ssl. This is setup *before* the nginx server is started.

```
$ kubectl delete deployments,svc my-nginx; kubectl create -f ./nginx-secure-app.yaml
```

At this point you can reach the nginx server from any node.

```
$ kubectl get pods -o yaml | grep -i podip
podIP: 10.244.3.5
node $ curl -k https://10.244.3.5
...
<h1>Welcome to nginx!</h1>
```

Note how we supplied the -k parameter to curl in the last step, this is because we don't know anything about the pods running nginx at certificate generation time, so we have to tell curl to ignore the CName mismatch. By creating a Service we linked the CName used in the certificate with the actual DNS name used by pods during Service lookup. Let's test this from a pod (the same secret is being reused for simplicity, the pod only needs nginx.crt to access the Service):

[service/networking/curlpod.yaml](#)

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: curl-deploymentspec:
  selector:
    matchLabels:
      app: curlpod
  replicas: 1
  template:
    metadata:
      labels:
        app: curlpod
    spec:
      volumes:
        - name: secret-volume
          secret:
            secretName: nginxsecret
      containers:
        - name: curlpod
          command:
            - sh
            - -c
            - while true; do sleep 1; done
          image: radial/busyboxplus:curl
          volumeMounts:
            - mountPath: /etc/nginx/ssl
              name: secret-volume

```

```

$ kubectl create -f ./curlpod.yaml
$ kubectl get pods -l app=curlpod
NAME                                READY  STATUS   RESTARTS  AGE
curl-deployment-1515033274-1410r    1/1    Running  0         1m
$ kubectl exec curl-deployment-1515033274-1410r -- curl https://my-nginx --cacert
/etc/nginx/ssl/nginx.crt
...
<title>Welcome to nginx!</title>
...

```

Exposing the Service

For some parts of your applications you may want to expose a Service onto an external IP address. Kubernetes supports two ways of doing this: NodePorts and LoadBalancers. The

Service created in the last section already used NodePort, so your nginx HTTPS replica is ready to serve traffic on the internet if your node has a public IP.

```
$ kubectl get svc my-nginx -o yaml | grep nodePort -C 5
```

```
uid: 07191fb3-f61a-11e5-8ae5-42010af00002
```

```
spec:
```

```
clusterIP: 10.0.162.149
```

```
ports:
```

```
- name: http
```

```
nodePort: 31704
```

```
port: 8080
```

```
protocol: TCP
```

```
targetPort: 80
```

```
- name: https
```

```
nodePort: 32453
```

```
port: 443
```

```
protocol: TCP
```

```
targetPort: 443
```

```
selector:
```

```
run: my-nginx
```

```
$ kubectl get nodes -o yaml | grep ExternalIP -C 1
```

```
- address: 104.197.41.11
```

```
type: ExternalIP
```

```
allocatable:
```

```
--
```

```
- address: 23.251.152.56
```

```
type: ExternalIP
```

```
allocatable:
```

```
...
```

```
$ curl https://<EXTERNAL-IP>:<NODE-PORT> -k
```

```
...
```

```
<h1>Welcome to nginx!</h1>
```

Let's now recreate the Service to use a cloud load balancer, just change the Type of my-nginx Service from NodePort to LoadBalancer:

```
$ kubectl edit svc my-nginx
```

```
$ kubectl get svc my-nginx
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
my-nginx	ClusterIP	10.0.162.149	162.222.184.144	80/TCP,81/TCP,82/TCP	21s

```
$ curl https://<EXTERNAL-IP> -k
```

...

<title>Welcome to nginx!</title>

The IP address in the EXTERNAL-IP column is the one that is available on the public internet. The CLUSTER-IP is only available inside your cluster/private cloud network.

Note that on AWS, type LoadBalancer creates an ELB, which uses a (long) hostname, not an IP. It's too long to fit in the standard `kubectl get svc` output, in fact, so you'll need to do `kubectl describe service my-nginx` to see it. You'll see something like this:

\$ `kubectl describe service my-nginx`

...

LoadBalancer Ingress:

a320587ffd19711e5a37606cf4a74574-1142138393.us-east-1.elb.amazonaws.com

...