

Kubectl Context and Configuration

`kubectl config view` *# Show Merged kubeconfig settings.*

use multiple kubeconfig files at the same time and view merged config

`KUBECONFIG=~/.kube/config:~/.kube/kubconfig2 kubectl config view`

Get the password for the e2e user

`kubectl config view -o jsonpath='{.users[?(@.name == "e2e")].user.password}'`

`kubectl config current-context` *# Display the current-context*

`kubectl config use-context my-cluster-name` *# set the default context to my-cluster-name*

add a new cluster to your kubeconf that supports basic auth

`kubectl config set-credentials kubeuser/foo.kubernetes.com --username=kubeuser
--password=kubepassword`

set a context utilizing a specific username and namespace.

`kubectl config set-context gce --user=cluster-admin --namespace=foo \
&& kubectl config use-context gce`

Creating Objects

Kubernetes manifests can be defined in json or yaml. The file extension .yaml, .yml, and .json can be used.

`kubectl create -f ./my-manifest.yaml` *# create resource(s)*

`kubectl create -f ./my1.yaml -f ./my2.yaml` *# create from multiple files*

`kubectl create -f ./dir` *# create resource(s) in all manifest files in dir*

`kubectl create -f https://git.io/vPieo` *# create resource(s) from url*

`kubectl run nginx --image=nginx` *# start a single instance of nginx*

`kubectl explain pods,svc` *# get the documentation for pod and svc manifests*

Create multiple YAML objects from stdin

```
cat <<EOF | kubectl create -f -
apiVersion: v1
kind: Pod
metadata:
  name: busybox-sleep
spec:
  containers:
  - name: busybox
    image: busybox
    args:
    - sleep
    - "1000000"
```

```
apiVersion: v1
kind: Pod
metadata:
  name: busybox-sleep-less
spec:
  containers:
  - name: busybox
    image: busybox
    args:
    - sleep
    - "1000"
```

EOF

Create a secret with several keys

```
cat <<EOF | kubectl create -f -
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
type: Opaque
data:
  password: $(echo -n "s33msi4" | base64 -w0)
  username: $(echo -n "jane" | base64 -w0)
EOF
```

Viewing, Finding Resources

Get commands with basic output

```
kubectl get services          # List all services in the namespace
kubectl get pods --all-namespaces # List all pods in all namespaces
kubectl get pods -o wide      # List all pods in the namespace, with more details
kubectl get deployment my-dep # List a particular deployment
kubectl get pods --include-uninitialized # List all pods in the namespace, including
uninitialized ones
```

Describe commands with verbose output

```
kubectl describe nodes my-node
kubectl describe pods my-pod
```

```
kubectl get services --sort-by=.metadata.name # List Services Sorted by Name
```

List pods Sorted by Restart Count

```
kubectl get pods --sort-by='.status.containerStatuses[0].restartCount'
```

Get the version label of all pods with label app=cassandra

```
kubectl get pods --selector=app=cassandra rc -o \
  jsonpath='{.items[*].metadata.labels.version}'
```

Get all running pods in the namespace

```
kubectl get pods --field-selector=status.phase=Running
```

Get ExternalIPs of all nodes

```
kubectl get nodes -o jsonpath='{.items[*].status.addresses[?(@.type=="ExternalIP")].address}'
```

List Names of Pods that belong to Particular RC

"jq" command useful for transformations that are too complex for jsonpath, it can be found at <https://stedolan.github.io/jq/>

```
sel=${$(kubectl get rc my-rc --output=json | jq -j '.spec.selector | to_entries | .[] |
"\(.key)=\(.value),")"%?}
echo $(kubectl get pods --selector=$sel --output=jsonpath='{.items..metadata.name}')
```

Check which nodes are ready

```
JSONPATH='{range .items[*]}{@.metadata.name}:{range
@.status.conditions[*]}{@.type}={@.status};{end}{end}' \
&& kubectl get nodes -o jsonpath="$JSONPATH" | grep "Ready=True"
```

List all Secrets currently in use by a pod

```
kubectl get pods -o json | jq '.items[].spec.containers[].env[]?.valueFrom.secretKeyRef.name' |
```

```
grep -v null | sort | uniq
```

List Events sorted by timestamp

```
kubectl get events --sort-by=.metadata.creationTimestamp
```

Updating Resources

As of version 1.11 rolling-update have been deprecated (see [CHANGELOG-1.11.md](#)), use rollout instead.

```
kubectl set image deployment/frontend www=image:v2          # Rolling update "www"
containers of "frontend" deployment, updating the image
kubectl rollout undo deployment/frontend                    # Rollback to the previous deployment
kubectl rollout status -w deployment/frontend               # Watch rolling update status of
"frontend" deployment until completion
```

deprecated starting version 1.11

```
kubectl rolling-update frontend-v1 -f frontend-v2.json      # (deprecated) Rolling update pods
of frontend-v1
kubectl rolling-update frontend-v1 frontend-v2 --image=image:v2 # (deprecated) Change the
name of the resource and update the image
kubectl rolling-update frontend --image=image:v2            # (deprecated) Update the pods
image of frontend
kubectl rolling-update frontend-v1 frontend-v2 --rollback   # (deprecated) Abort existing
rollout in progress
```

```
cat pod.json | kubectl replace -f -                          # Replace a pod based on the JSON passed
into std
```

Force replace, delete and then re-create the resource. Will cause a service outage.

```
kubectl replace --force -f ./pod.json
```

Create a service for a replicated nginx, which serves on port 80 and connects to the containers on port 8000

```
kubectl expose rc nginx --port=80 --target-port=8000
```

Update a single-container pod's image version (tag) to v4

```
kubectl get pod mypod -o yaml | sed 's/(image: myimage\):.*$/\1:v4/' | kubectl replace -f -
```

```
kubectl label pods my-pod new-label=awesome                # Add a Label
```

```
kubectl annotate pods my-pod icon-url=http://goo.gl/XXBTWq  # Add an annotation
```

```
kubectl autoscale deployment foo --min=2 --max=10          # Auto scale a deployment "foo"
```

Patching Resources

```
kubectl patch node k8s-node-1 -p '{"spec":{"unschedulable":true}}' # Partially update a node
```

Add a new element to a positional array

```
kubectl patch sa default --type='json' -p='[{"op": "add", "path": "/secrets/1", "value": {"name": "whatever" } }]'
```

Update a container's image; spec.containers[].name is required because it's a merge key*

```
kubectl patch pod valid-pod -p '{"spec":{"containers":[{"name":"kubernetes-serve-hostname","image":"new image"}]}}'
```

Update a container's image using a json patch with positional arrays

```
kubectl patch pod valid-pod --type='json' -p='[{"op": "replace", "path": "/spec/containers/0/image", "value":"new image"}]'
```

Disable a deployment livenessProbe using a json patch with positional arrays

```
kubectl patch deployment valid-deployment --type json -p='[{"op": "remove", "path": "/spec/template/spec/containers/0/livenessProbe"}]'
```

Editing Resources

The edit any API resource in an editor.

```
kubectl edit svc/docker-registry # Edit the service named docker-registry
```

```
KUBE_EDITOR="nano" kubectl edit svc/docker-registry # Use an alternative editor
```

Scaling Resources

```
kubectl scale --replicas=3 rs/foo # Scale a replicaset named 'foo' to 3
```

```
kubectl scale --replicas=3 -f foo.yaml # Scale a resource specified in "foo.yaml" to 3
```

```
kubectl scale --current-replicas=2 --replicas=3 deployment/mysql # If the deployment named mysql's current size is 2, scale mysql to 3
```

```
kubectl scale --replicas=5 rc/foo rc/bar rc/baz # Scale multiple replication controllers
```

Deleting Resources

```
kubectl delete -f ./pod.json # Delete a pod using the type and name specified in pod.json
```

```

kubectrl delete pod,service baz foo # Delete pods and services with same names "baz" and "foo"
kubectrl delete pods,services -l name=myLabel # Delete pods and services with label name=myLabel
kubectrl delete pods,services -l name=myLabel --include-uninitialized # Delete pods and services, including uninitialized ones, with label name=myLabel
kubectrl -n my-ns delete po,svc --all # Delete all pods and services, including uninitialized ones, in namespace my-ns,

```

Interacting with running Pods

```

kubectrl logs my-pod # dump pod logs (stdout)
kubectrl logs my-pod --previous # dump pod logs (stdout) for a previous instantiation of a container
kubectrl logs my-pod -c my-container # dump pod container logs (stdout, multi-container case)
kubectrl logs my-pod -c my-container --previous # dump pod container logs (stdout, multi-container case) for a previous instantiation of a container
kubectrl logs -f my-pod # stream pod logs (stdout)
kubectrl logs -f my-pod -c my-container # stream pod container logs (stdout, multi-container case)
kubectrl run -i --tty busybox --image=busybox -- sh # Run pod as interactive shell
kubectrl attach my-pod -i # Attach to Running Container
kubectrl port-forward my-pod 5000:6000 # Listen on port 5000 on the local machine and forward to port 6000 on my-pod
kubectrl exec my-pod -- ls / # Run command in existing pod (1 container case)
kubectrl exec my-pod -c my-container -- ls / # Run command in existing pod (multi-container case)
kubectrl top pod POD_NAME --containers # Show metrics for a given pod and its containers

```

Interacting with Nodes and Cluster

```

kubectrl cordon my-node # Mark my-node as unschedulable
kubectrl drain my-node # Drain my-node in preparation for maintenance
kubectrl uncordon my-node # Mark my-node as schedulable
kubectrl top node my-node # Show metrics for a given node
kubectrl cluster-info # Display addresses of the master and services
kubectrl cluster-info dump # Dump current cluster state to stdout
kubectrl cluster-info dump --output-directory=/path/to/cluster-state # Dump current cluster state to /path/to/cluster-state

```

If a taint with that key and effect already exists, its value is replaced as specified.

kubectl taint nodes foo **dedicated**=special-user:NoSchedule

Resource types

kubectl api-resources --namespaced=**true** *# All namespaced resources*

kubectl api-resources --namespaced=**false** *# All non-namespaced resources*

kubectl api-resources -o name *# All resources with simple output (just the resource name)*

kubectl api-resources -o wide *# All resources with expanded (aka "wide") output*

kubectl api-resources --verbs=list,get *# All resources that support the "list" and "get" request verbs*

kubectl api-resources --api-group=extensions *# All resources in the "extensions" API group*

Formatting output

To output details to your terminal window in a specific format, you can add either the -o or --output flags to a supported kubectl command.

Output format	Description
-o=custom-columns=<spec>	Print a table using a comma separated list of custom columns
-o=custom-columns-file=<filename>	Print a table using the custom columns template in the <filename> file
-o=json	Output a JSON formatted API object
-o=jsonpath=<template>	Print the fields defined in a jsonpath expression
-o=jsonpath-file=<filename>	Print the fields defined by the jsonpath expression in the <filename> file
-o=name	Print only the resource name and nothing else
-o=wide	Output in the plain-text format with any additional information, and for pods, the node name is included
-o=yaml	Output a YAML formatted API object

Kubectl output verbosity and debugging

Kubectl verbosity is controlled with the `-v` or `--v` flags followed by an integer representing the log level. General Kubernetes logging conventions and the associated log levels are described [here](#).

Verbosity	Description
- -v=0	Generally useful for this to ALWAYS be visible to an operator.
- -v=1	A reasonable default log level if you don't want verbosity.
- -v=2	Useful steady state information about the service and important log messages that may correlate to significant changes in the system. This is the recommended default log level for most systems.
- -v=3	Extended information about changes.
- -v=4	Debug level verbosity.
- -v=6	Display requested resources.
- -v=7	Display HTTP request headers.
- -v=8	Display HTTP request contents.
- -v=9	Display HTTP request contents without truncation of contents.